

# Parallel Matrix Multiplication Study

John Nehls

10/11/15

**Note:** This project may be found on GitHub at <https://github.com/JohnNehls/ParallelMatrixMultiplication>

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Parallel Matrix Multiplication</b>	<b>2</b>
<b>3</b>	<b>Results</b>	<b>3</b>
3.1	Details of the Experiments . . . . .	3
3.2	Total Matrix Calculation Time . . . . .	3
3.2.1	Results . . . . .	3
3.3	Ratio of Calculation Time Over Communication Time . . . . .	5
	<b>Appendices</b>	<b>7</b>
<b>A</b>	<b>Serial Code</b>	<b>7</b>
<b>B</b>	<b>OpenMP Code</b>	<b>9</b>
<b>C</b>	<b>MPI Code</b>	<b>11</b>
<b>D</b>	<b>Matrix.h</b>	<b>14</b>

# 1 Introduction

This document will serve as a comparison of serial, OpenMP parallel, and MPI parallel code that accomplishes the same task: matrix multiplication. Along with comparing the total matrix multiplication times of the codes, we will look at the ratio of time spent calculating the multiplication to the time the parallel tool spends communicating data. Then, we attempt to explain why the parallelization schemes scale the way they do.

## 2 Parallel Matrix Multiplication

The scheme used to distribute the matrix multiplication of matrices A and B

$$AB = C \tag{1}$$

is the same approach as the one displayed in Lecture 2 of the course. Thus, to offer a brief summary of the approach, we will use slightly modified figures of ones shown in class (which seem to be from an excellent introduction to parallel computing [1])

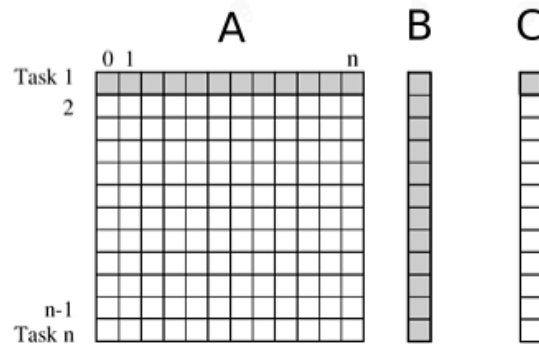


Figure 1:

In figure 1 we see how the matrix multiplication is split across n tasks. Task 1 computes the the first element of C using the information in dark gray: the first row of A and all of matrix B. This is a general way to see the how the computation may be split in to n tasks. In practice often each task will calculate more than one quantity, this is the concept of granularity. In figure 2 we see how the multiplication of equation 1 is split over 4 tasks.

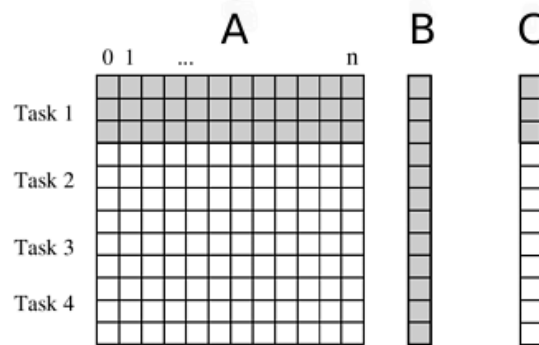


Figure 2:

The parallelization scheme used in our MPI and OpenMP codes is the same as shown here: each task sees one or more rows of A, all of B, and calculates a subset of C. For a more detailed treatment, I encourage you to check out the book or visit the website hyper linked [1].

## 3 Results

### 3.1 Details of the Experiments

To carry out the comparisons of serial to OpenMP to MPI, we calculated the multiplication of random-valued, square matrices of dimensions 240x240 and 2400x2400. The parallel codes computed these matrices using 2, 4, 8, 16, 32, 64, and 128 parallel processes. Each result was averaged over 10 runs. The Arizona Center of Mathematical Sciences at the University of Arizona provided me with time on a Silocon Graphics (SGI) UV2000 which contains, after recent upgrades, 128 Intel Xeon E5-4617 "Sandy Bridge" 6 core Processors at 2.9GHz[2].

### 3.2 Total Matrix Calculation Time

First, to understand the total matrix multiplication timing data, we look at the postcode for the serial, OpenMP, and MPI total matrix multiplication times. The important detail in this code is that the creation and filling of the matrices is excluded from the timing.

```
1 int main(int argc, char *argv[]) {
2   A,B = CreateMatrices();           /* Excluded from timing!!! */
3   FillMatricesRandomly(A, B);       /* Excluded from timing!!! */
4
5   start = currentTime()
6
7   // Matrix Multiplication
8
9   end = currentTime()
10  double matrixCalculationTime = end - start;
11
12  return 0;
13 }
```

#### 3.2.1 Results

To get the point, lets look at the total multiplication time results.

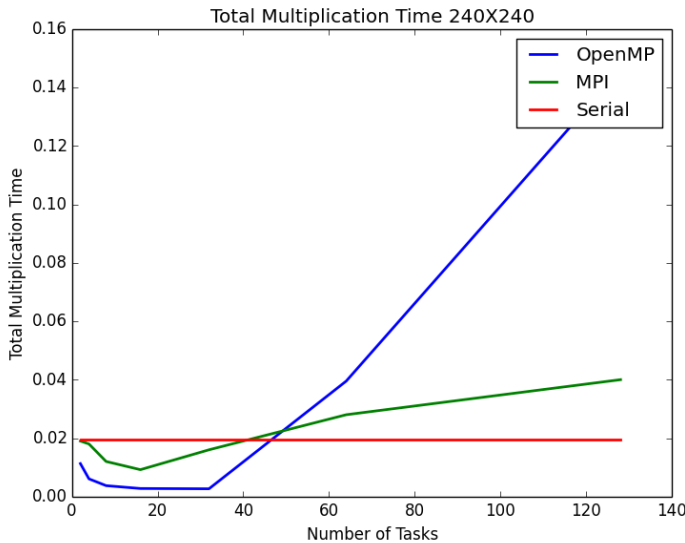


Figure 3:

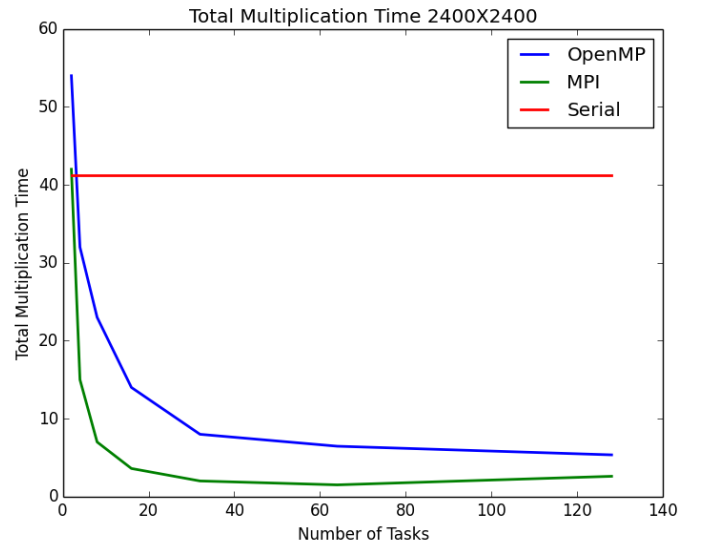


Figure 4:

Now to explain the behavior, we will use Amdahl's Law. Amdahl's Law states

$$speedup = 1/(P/N + S) \quad (2)$$

where P = parallel fraction, N = number of processors and S = serial fraction [3]. Extending this a little further, we find the total multiplication time

$$totaltime = time0/speedup = time0 \times (P/N + S), \quad (3)$$

where  $time0$  is the amount of time it takes for the serial code to execute the task at hand.

To explain the dependence of total matrix multiplication time on the number of tasks, we plot Amdahl's Law in the form shown in equation 3 with a parallel fraction of 1 in figures 5 and 6, showing the ideal case, which, it should be said, should be the correct parallel fraction since no serial actions are timed.

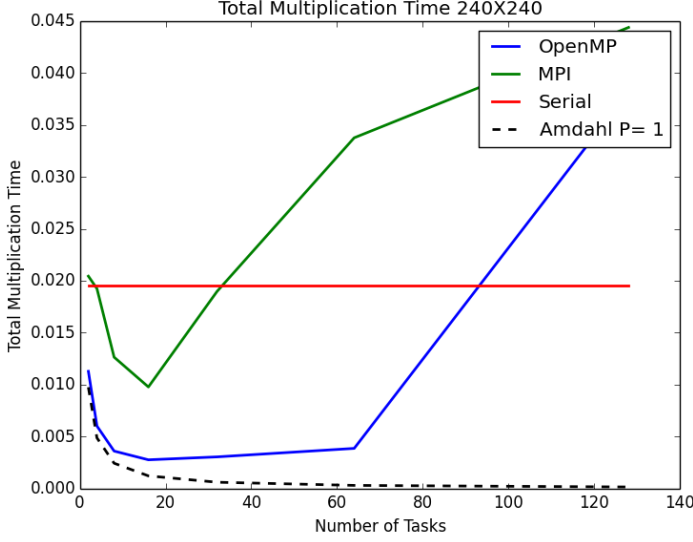


Figure 5:

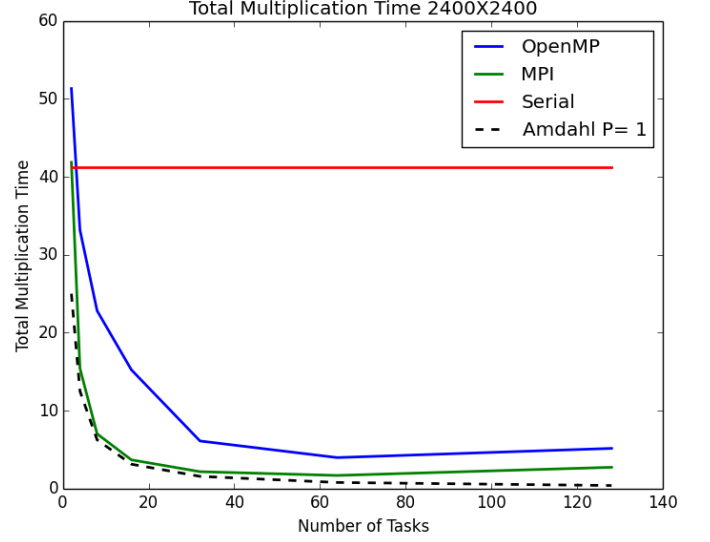


Figure 6:

First of all we see that in figure 5 the experiment is very different from the theory, not even displaying the same shape, where as, for computation of the larger matrices, figure 6, the theory predicts the behavior of the experiment well. The simple explanation of this is that the slaves( or threads) in the 240x240 were spending more time communicating than computing, which, when brought to the extreme, here a granularity of about 50 rows, the computation begins to take longer in parallel than it did in for the serial code. For the 2400x2400 case, this extreme saturation did not appear. We will visit this again after we look at the ratio of computation over communication times in the next results section, which should confirm or deny this explanation for the far from theoretical result in figure 5.

### 3.3 Ratio of Calculation Time Over Communication Time

Here we calculate the ratio of the computation time over the communication time by assuming

$$T_{total} = T_{comp} + T_{comm}, \quad (4)$$

where  $T_{total}$  is the total matrix multiplication time,  $T_{comp}$  is the computation time, and  $T_{comm}$  is the communication time of the parallel tool, to be true. If so, we can say that we can find the communication time by

$$T_{comm} = T_{total} - T_{comp}. \quad (5)$$

So how do we acquire the computation time? Here we provide the pseudo code for acquiring the computation time from the MPI code.

```
1  int main(int argc, char *argv[]) {
2      /* initialize MPI */
3      A,B = CreateMatrices();          /* Excluded from timing!!! */
4
5      if (rank == 0) { /* Master initializes work*/
6          FillMatricesRandomly( A, B);    /* Excluded from timing!!! */
7
8          //send data to each slave needs to calculate its part of the workload
9          MPI_Isend(&informationToSlave);
10     }
11
12     if (rank > 0) { /* work done by slaves (not rank = 0)*/
13         //receive data from master to calculate part of the workload
14         MPI_Recv(&information);
15
16         //FOR THE COMPUTATIONAL TIME: start time for local time: the amount of time to do matrix
17         //calculation for this process
18         localTimeSaver = MPI_Wtime();
19
20         /* DO Matrix Multiplication-- Loops that do the computation */
21
22         //FOR THE COMPUTATIONAL TIME: calculate local time: the amount of time to do matrix
23         //calculation for this process
24         localTimeSaver = MPI_Wtime() - localTimeSaver; // calculates the time spent on calculation
25
26         //send back the matrix calc data to master
27         MPI_Isend(&calculatedMatrixInfo);
28
29         //FOR THE COMPUTATIONAL TIME: localTimeSaver to master
30         MPI_Isend(&localTimeSaver);
31     }
32     if (rank == 0) { /* master gathers processed work*/
33
34         //receive matrix data from slaves
35         MPI_Recv(&CalculatedMatrixInformation);
36
37         //FOR THE COMPUTATIONAL TIME
38         // find the longest local calculation (which we take as the total amount of calculation time,
39         // the rest coming from communication.
40         MPI_Recv(&localComputationTimeForEachSlave);
41         double maxLocalMultiplicationTime = Max(localComputationTimeForEachSlave)
42     }
43
44     return 0;
45 }
```

Here we provide the pseudo code for acquiring the average computation time from the OpenMP code. I am unsure if this is an accurate measure, but it is my best guess on how to do the timing, so I decided to in this report. Note, again, that the average computation time is found, not the computation time.

```

1  int main(int argc, char *argv[]) {
2      A,B = CreateMatrices();           /* Excluded from timing!!! */
3      FillMatricesRandomly(A, B);       /* Excluded from timing!!! */
4
5      omp_set_num_threads(numThreads); // set the number of threads
6
7      // Matrix Multiplication
8      #pragma omp parallel for private(val) reduction(+:sum)
9      for (int i = 0; i < A.rows(); i++) { //iterate through rows of A (parallelized loop)
10
11          val = omp_get_wtime(); // for each thread get Start time
12
13          for (int j = 0; j < B.cols(); j++) { //iterate through columns of B
14              for (int k = 0; k < B.rows(); k++) { //iterate through rows of B
15                  C(i,j) += (A(i,k) * B(k,j));
16              }
17          }
18          sum += omp_get_wtime() - val; //add up ALL OF THE COMPUTATION TIMES
19      }
20      AverageComputationTime = sum/numberOfThreads
21
22      return 0;
23 }

```

Now that we have the the ratio of the computation time over the communication time (comp/comm), we can look at the results in figures 7 and 8 and see that our intuition was true from the previous results section when trying to describe why the the total multiplication time was so large for the 240x240 parallel runs. The (comp/comm) becomes much less than one for both MPI and OpenMP for the 240x240, and, in the 2400x2400 case, only just approaches one, keeping it near Amdahl's Law equation 2.

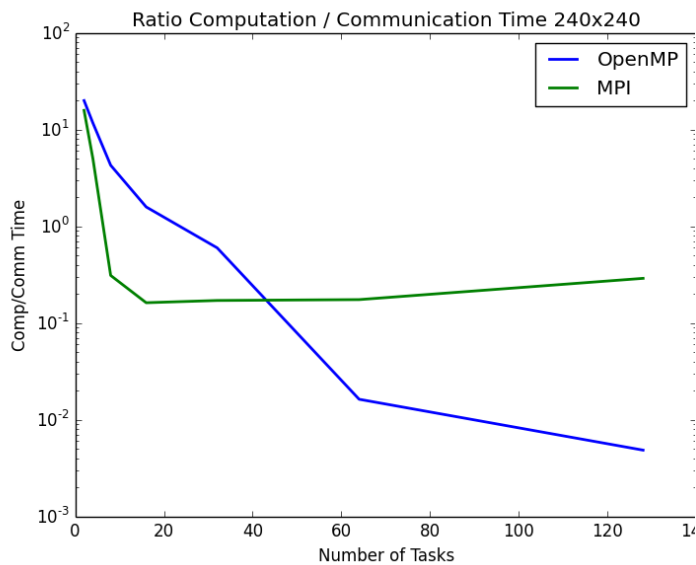


Figure 7:

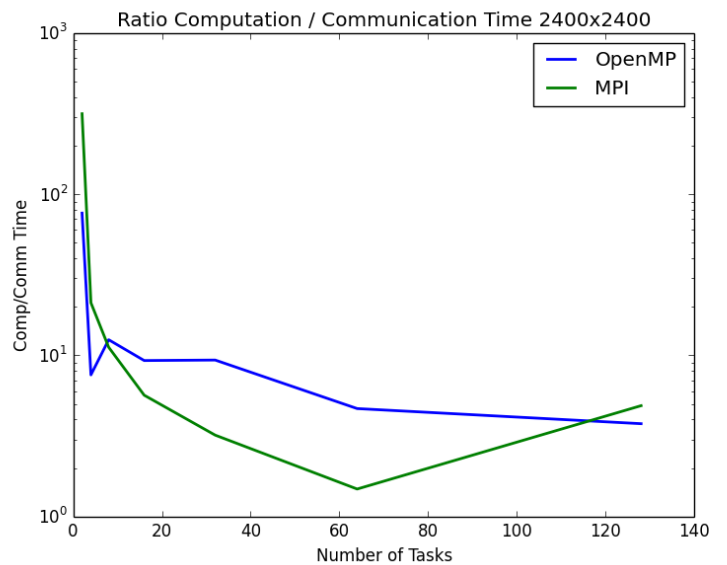


Figure 8:

# Appendices

## A Serial Code

```
1 #include<stdio.h>
2 #include<cstdlib>
3 #include<time.h>
4 #include<iostream>
5 #include<matrix.h>
6
7 /* Function Prototypes */
8 void FillMatricesRandomly(Matrix<double> &A, Matrix<double> &B );
9 void PrintMatrices( Matrix<double> &A, Matrix<double> &B, Matrix<double> &C );
10
11 /* Global variables that could be inputs #TODO */
12 int randomHigh = 100;          /* the upper bound of the random numbers that fill A and B*/
13 int randomLow = 0;             /* the lower bound of the random numbers that fill A and B*/
14
15 int main(int argc, char *argv[]) {
16     std::cout << "Starting a serial matrix multiplication. \n " << std::endl;
17     // Read in the input: N
18     if ( argv[1]== NULL ){ // check input is supplied
19         std::cout << "ERROR: The program must be executed in the following way \n\n \t \"./serial.
20             exe N \" \n\n where N is an integer. \n \n \" << std::endl;
21         return 1;
22     }
23     int N = atoi(argv[1]); // The dimensions of the matrices MUST be specified at runtime.
24     std::cout << "The matrices are: " << N<<"x"<<N<< std::endl;
25     // for simplicity, all matrices will be N x N
26     int numberOfRowsA = N; int numberOfColsA = N; int numberOfRowsB = N; int numberOfColsB = N;
27
28     //Declare matrices: Matix class is a 2D vector
29     Matrix<double> A = Matrix<double>(numberOfRowsA, numberOfColsA);
30     Matrix<double> B = Matrix<double>(numberOfRowsB, numberOfColsB);
31     Matrix<double> C = Matrix<double>(numberOfRowsA, numberOfColsB);
32
33     FillMatricesRandomly(A, B); /* Excluded from timing!!! */
34
35     struct timespec start, end;
36     clock_gettime(CLOCK_MONOTONIC, &start); // start timing
37
38     // Matrix Multiplication
39     for (int i = 0; i < A.rows(); i++) { //iterate through rows of A
40         for (int j = 0; j < B.cols(); j++) { //iterate through columns of B
41             for (int k = 0; k < B.rows(); k++) { //iterate through rows of B
42                 C(i,j) += (A(i,k) * B(k,j));
43             }
44         }
45     }
46
47     clock_gettime(CLOCK_MONOTONIC, &end); // end timing
48     double matrixCalculationTime = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec)*1e
49         -9;
50     std::cout << "\nTotal multiplication time = " << matrixCalculationTime << std::endl;
51     // PrintMatrices(A, B, C);
52     return 0;
53 }
54
55 void FillMatricesRandomly(Matrix<double> &A, Matrix<double> &B){
56     /* initialize the random number generator with the current time */
57     srand( time( NULL ) );
58     for (int i = 0; i < A.rows(); i++) {
59         for (int j = 0; j < A.cols(); j++) {
60             A(i,j) = rand() % (randomHigh - randomLow) + randomLow;
61         }
62     }
63     for (int i = 0; i < B.rows(); i++) {
64         for (int j = 0; j < B.cols(); j++) {
65             B(i,j) = rand() % (randomHigh - randomLow) + randomLow;
66         }
67     }
68 }
```

```
67
68
69 void PrintMatrices(Matrix<double> &A, Matrix<double> &B, Matrix<double> &C){
70     for (int i = 0; i < A.rows(); i++) {
71         std::cout << "\n" << std::endl;
72         for (int j = 0; j < A.cols(); j++)
73             std::cout << A(i,j) << " ";
74     }
75     std::cout << "\n\n" << std::endl;
76     for (int i = 0; i < B.rows(); i++) {
77         std::cout << "\n" << std::endl;
78         for (int j = 0; j < B.cols(); j++)
79             std::cout << B(i,j) << " ";
80     }
81     std::cout << "\n\n" << std::endl;
82     for (int i = 0; i < C.rows(); i++) {
83         std::cout << "\n" << std::endl;
84         for (int j = 0; j < C.cols(); j++)
85             std::cout << C(i,j) << " ";
86     }
87     std::cout << "\n\n" << std::endl;
88 }
```



## B OpenMP Code

```
1 #include<cstdlib>
2 #include<time.h>
3 #include<iostream>
4 #include<matrix.h>
5 #include<omp.h>
6
7
8 /* Function Prototypes */
9 void FillMatricesRandomly(Matrix<double> &A, Matrix<double> &B );
10 void PrintMatrices( Matrix<double> &A, Matrix<double> &B, Matrix<double> &C );
11
12 /* Global variables that could be inputs #TODO */
13 int randomHigh = 100;          /* the upper bound of the random numbers that fill A and B*/
14 int randomLow = 0;             /* the lower bound of the random numbers that fill A and B*/
15
16 int main(int argc, char *argv[]) {
17     std::cout << "Starting an OpenMP parallel matrix multiplication. \n " << std::endl;
18     // Read in the two inputs: NumberOfOMPthreads and N
19     if ( argv[1]== NULL || argv[2] == NULL){ // check if inputs were supplied
20         std::cout << "ERROR: The program must be executed in the following way \n\n \t \"./omp.exe
21             NumberOfThreads N \" \n\n where NuberOfThreads and N are integers. \n \n " << std::endl;
22         return 1;
23     }
24     int numThreads = atoi(argv[1]);
25     std::cout << "The number of OpenMP threads: " << numThreads << std::endl;
26     omp_set_dynamic(0);          // do not allow the number of threads to be set internally
27     omp_set_num_threads(numThreads); // set the number of threads
28
29     int N = atoi(argv[2]); // The dimensions of the matrices MUST be specified at runtime.
30     std::cout << "The matrices are: " << N<<"x"<<N<< std::endl;
31     // for simplicity, all matrices will be N×N
32     int numberOfRowsA = N; int numberOfColsA = N; int numberOfRowsB = N; int numberOfColsB = N;
33     //Declare matrices: Matix class is a 2D vector
34     Matrix<double> A = Matrix<double>(numberOfRowsA, numberOfColsA);
35     Matrix<double> B = Matrix<double>(numberOfRowsB, numberOfColsB);
36     Matrix<double> C = Matrix<double>(numberOfRowsA, numberOfColsB);
37
38     FillMatricesRandomly(A, B); /* Excluded from timing!!! */
39
40     struct timespec start, end;
41     clock_gettime(CLOCK_MONOTONIC, &start); // start timing
42
43     // Used to calculate the longest a process spends calculating its part of the workload.
44     // double sumLocalTime = 0;
45     double sum = 0;
46     double val = 0;
47     // Matrix Multiplication
48     #pragma omp parallel for private(val) reduction(+:sum)
49     for (int i = 0; i < A.rows(); i++) { //iterate through rows of A (parallelized loop)
50         val = omp_get_wtime();
51         for (int j = 0; j < B.cols(); j++) { //iterate through columns of B
52             for (int k = 0; k < B.rows(); k++) { //iterate through rows of B
53                 C(i,j) += (A(i,k) * B(k,j));
54             }
55         }
56         sum += omp_get_wtime() - val;
57     }
58
59     clock_gettime(CLOCK_MONOTONIC, &end); // end timing
60     double totalMatrixCalculationTime = (end.tv_sec - start.tv_sec) + (end.tv_nsec - start.tv_nsec)
61         *1e-9;
62     std::cout << "Total multiplication time = " << totalMatrixCalculationTime << std::endl;
63     std::cout << "average multiplication time = " << sum/numThreads << std::endl;
64     std::cout << "Approximate Communication time = " << totalMatrixCalculationTime - sum/numThreads
65         << std::endl;
66     // PrintMatrices(A, B, C);
67     return 0;
68 }
69
70 void FillMatricesRandomly(Matrix<double> &A, Matrix<double> &B){
71     /* initialize the random number generator with the current time */
```

```

69 | srand( time( NULL ) );
70 | for (int i = 0; i < A.rows(); i++) {
71 |     for (int j = 0; j < A.cols(); j++) {
72 |         A(i,j) = rand() % (randomHigh - randomLow) + randomLow;
73 |     }
74 | }
75 | for (int i = 0; i < B.rows(); i++) {
76 |     for (int j = 0; j < B.cols(); j++) {
77 |         B(i,j) = rand() % (randomHigh - randomLow) + randomLow;
78 |     }
79 | }
80 | }
81 |
82 | void PrintMatrices(Matrix<double> &A, Matrix<double> &B, Matrix<double> &C){
83 |     for (int i = 0; i < A.rows(); i++) {
84 |         std::cout << "\n" << std::endl;
85 |         for (int j = 0; j < A.cols(); j++)
86 |             std::cout << A(i,j) << " ";
87 |     }
88 |     std::cout << "\n\n" << std::endl;
89 |     for (int i = 0; i < B.rows(); i++) {
90 |         std::cout << "\n" << std::endl;
91 |         for (int j = 0; j < B.cols(); j++)
92 |             std::cout << B(i,j) << " ";
93 |     }
94 |     std::cout << "\n\n" << std::endl;
95 |     for (int i = 0; i < C.rows(); i++) {
96 |         std::cout << "\n" << std::endl;
97 |         for (int j = 0; j < C.cols(); j++)
98 |             std::cout << C(i,j) << " ";
99 |     }
100 |     std::cout << "\n\n" << std::endl;
101 | }
102 | }

```

## C MPI Code

```
1 #include<cstdlib>
2 #include<ctime>
3 #include<mpi.h>
4 #include<algorithm>
5 #include<matrix.h>
6
7 // Function Prototypes
8 void FillMatricesRandomly(Matrix<double> &A, Matrix<double> &B );
9 void PrintMatrices( Matrix<double> &A, Matrix<double> &B, Matrix<double> &C );
10
11 /* MPI Send and Recieve Tags */
12 #define ROW_START_TAG 0 //tag for communicating the start row of the workload for a slave
13 #define ROW_END_TAG 1 //tag for communicating the end row of the workload for a slave
14 #define A_ROWS_TAG 2 //tag for communicating the address of the data to be worked on to slave
15 #define C_ROWS_TAG 3 //tag for communicating the address of the calculated data to master
16 #define LOCAL_TIME_TAG 4//tag for communicating the address of the local matrix calculation time
    to master
17
18 // Instantiate global variables used in the parallelization
19 int rank; // mpi: process id number
20 int nProcesses; // mpi: number of total processes
21 MPI_Status status; // mpi: store status of a MPI_Recv
22 MPI_Request request; // mpi: capture request of a MPI_Isend
23 int rowStart, rowEnd; // which rows of A that are calculated by the slave process
24 int granularity; // granularity of parallelization (# of rows per processor)
25
26 //Used to calculate totalmultiplication time: communication + calculations
27 double start_time, end_time;
28 // Used to calculate the longest a process spends calculating its part of the workload.
29 double localTimeSaver;
30
31 /* Global variables that could be inputs #TODO */
32 int randomHigh = 100; // the upper bound of the random numbers that fill A and B
33 int randomLow = 0; // the lower bound of the random numbers that fill A and B
34
35 int main(int argc, char *argv[]) {
36 // ***** Handle the input size of the Matrices, N, where matrices A, B, and C will be NxN *****
37 if ( argv[1]== NULL ){// check if the input was supplied
38     std::cout << "ERROR: The program must be executed in the following way \n\n \t \"mpirun -n
        NumberOfProcesses mpi.exe N \" \n\n where N is an integer. \n \n \" << std::endl;
39     return 1;
40 }
41 int N = atoi(argv[1]); // The dimensions of the matrices MUST be specified at runtime.
42 // for simplicity, all matrices will be NxN
43 int numberOfRowsA = N; int numberOfColsA = N; int numberOfRowsB = N; int numberOfColsB = N;
44 //Declare matrices: Matrix class is a 2D vector
45 Matrix<double> A = Matrix<double>(numberOfRowsA, numberOfColsA);
46 Matrix<double> B = Matrix<double>(numberOfRowsB, numberOfColsB);
47 Matrix<double> C = Matrix<double>(numberOfRowsA, numberOfColsB);
48
49 // MPI:
50 MPI_Init(&argc, &argv); // initialize MPI */
51 MPI_Comm_rank(MPI_COMM_WORLD, &rank); // store the rank */
52 MPI_Comm_size(MPI_COMM_WORLD, &nProcesses); // store the number of processes */
53
54 if (rank == 0) { /* Master initializes work*/
55     std::cout << "Starting an MPI parallel matrix multiplication. \n " << std::endl;
56     std::cout << "The matrices are: " << N<<"x"<<N<< std::endl;
57     FillMatricesRandomly( A, B); // Excluded from timing!!! */
58
59     /* Begin Timing: used for total multiplication time: communication + calculations */
60     start_time = MPI_Wtime();
61     for (int i = 1; i < nProcesses; i++) { /* for each slave */
62         // calculate granularity (-1 comes from excluding the master process)
63         granularity = (numberOfRowsA / (nProcesses - 1));
64         rowStart = (i - 1) * granularity;
65         if (((i + 1) == nProcesses) && ((numberOfRowsA % (nProcesses - 1)) != 0)) { //if rows of [A]
            cannot be equally divided among slaves
66             rowEnd = numberOfRowsA; //last slave gets all the remaining rows
67         } else {
68             rowEnd = rowStart + granularity; //rows of [A] are equally divisable among slaves
```

```

69     }
70     //send the low bound, without blocking, to the intended slave
71     MPI_Isend(&rowStart, 1, MPI_INT, i, ROW_END_TAG, MPI_COMM_WORLD, &request);
72     //next send the upper bound without blocking, to the intended slave
73     MPI_Isend(&rowEnd, 1, MPI_INT, i, ROW_START_TAG, MPI_COMM_WORLD, &request);
74     //finally send the allocated row granularity of [A] without blocking, to the intended slave
75     MPI_Isend(&A(rowStart,0), (rowEnd - rowStart) * numberOfColsA, MPI_DOUBLE, i, A_ROWS_TAG,
76               MPI_COMM_WORLD, &request);
77 }
78 //broadcast B (MPI_Bcast: Broadcasts a message from the process with rank "root" to all other
79 //processes of the communicator)
80 MPI_Bcast(&B(0,0), numberOfRowsB*numberOfColsB, MPI_DOUBLE, 0, MPI_COMM_WORLD);
81 if (rank > 0) { /* work done by slaves (not rank = 0)*/
82     //receive low bound from the master
83     MPI_Recv(&rowStart, 1, MPI_INT, 0, ROW_END_TAG, MPI_COMM_WORLD, &status);
84     //next receive upper bound from the master
85     MPI_Recv(&rowEnd, 1, MPI_INT, 0, ROW_START_TAG, MPI_COMM_WORLD, &status);
86     //finally receive row granularity of [A] to be processed from the master
87     MPI_Recv(&A(rowStart,0), (rowEnd - rowStart) * numberOfColsA, MPI_DOUBLE, 0, A_ROWS_TAG,
88               MPI_COMM_WORLD, &status);
89     // start time for local time: the amount of time to do matrix calculation for this process
90     localTimeSaver = MPI_Wtime();
91     /* Matrix Multiplication */
92     for (int i = rowStart; i < rowEnd; i++) { //the given set of rows of A (parallelized loop)
93         for (int j = 0; j < B.cols(); j++) { //iterate through columns of [B]
94             for (int k = 0; k < B.rows(); k++) { //iterate through rows of [B]
95                 C(i,j) += (A(i,k) * B(k,j));
96             }
97         }
98     }
99     // calculate local time: the amount of time to do matrix calculation for this process
100    localTimeSaver = MPI_Wtime() - localTimeSaver;
101
102    //send back the low bound first without blocking, to the master
103    MPI_Isend(&rowStart, 1, MPI_INT, 0, ROW_END_TAG, MPI_COMM_WORLD, &request);
104    //send the upper bound next without blocking, to the master
105    MPI_Isend(&rowEnd, 1, MPI_INT, 0, ROW_START_TAG, MPI_COMM_WORLD, &request);
106    //finally send the processed granularity of data without blocking, to the master
107    MPI_Isend(&C(rowStart,0), (rowEnd - rowStart) * numberOfColsB, MPI_DOUBLE, 0, C_ROWS_TAG,
108              MPI_COMM_WORLD, &request);
109    //send back the local calculation time without blocking, to the master
110    MPI_Isend(&localTimeSaver, 1, MPI_DOUBLE, 0, LOCAL_TIME_TAG, MPI_COMM_WORLD, &request);
111
112 }
113
114 if (rank == 0) { /* master gathers processed work*/
115     for (int i = 1; i < nProcesses; i++) { // untill all slaves have handed back the processed
116         data
117         //receive low bound from a slave
118         MPI_Recv(&rowStart, 1, MPI_INT, i, ROW_END_TAG, MPI_COMM_WORLD, &status);
119         //receive upper bound from a slave
120         MPI_Recv(&rowEnd, 1, MPI_INT, i, ROW_START_TAG, MPI_COMM_WORLD, &status);
121         // //receive processed data from a slave
122         MPI_Recv(&C(rowStart,0), (rowEnd - rowStart) * numberOfColsB, MPI_DOUBLE, i, C_ROWS_TAG,
123                 MPI_COMM_WORLD, &status);
124     }
125     end_time = MPI_Wtime(); //end time of the total matrix matrix multiplication
126     double totalMultiplicationTime = end_time - start_time;
127
128     // find the longest local calculation (which we take as the total amount of calculation time,
129     // the rest coming from communication.
130     std::vector<double> LocalMultiplicationTimes = std::vector<double>(nProcesses);
131     for (int i = 1; i < nProcesses; i++) {
132         MPI_Recv(&LocalMultiplicationTimes[i], 1, MPI_DOUBLE, i, LOCAL_TIME_TAG, MPI_COMM_WORLD, &
133                 status);
134     }
135     double maxLocalMultiplicationTime = *std::max_element(LocalMultiplicationTimes.begin(),
136                   LocalMultiplicationTimes.end());

```

```

134 // print out the results
135 std::cout << "Total multiplication time = " << totalMultiplicationTime << "\n" << std::endl;
136 std::cout << "Longest multiplication time = " << maxLocalMultiplicationTime << "\n" << std::
    endl;
137 std::cout << "Approximate communication time = " << totalMultiplicationTime -
    maxLocalMultiplicationTime << "\n\n" << std::endl;
138
139 // PrintMatrices(A, B, C); // for debugging
140
141 }
142 MPI_Finalize(); //finalize MPI operations
143 return 0;
144 }
145
146 void FillMatricesRandomly(Matrix<double> &A, Matrix<double> &B){
147 /* initialize the random number generator with the current time */
148 srand( time( NULL ) );
149 for (int i = 0; i < A.rows(); i++) {
150     for (int j = 0; j < A.cols(); j++) {
151         A(i,j) = rand() % (randomHigh - randomLow) + randomLow;
152     }
153 }
154 for (int i = 0; i < B.rows(); i++) {
155     for (int j = 0; j < B.cols(); j++) {
156         B(i,j) = rand() % (randomHigh - randomLow) + randomLow;
157     }
158 }
159 }
160
161
162 void PrintMatrices(Matrix<double> &A, Matrix<double> &B, Matrix<double> &C){
163     for (int i = 0; i < A.rows(); i++) {
164         std::cout << "\n" << std::endl;
165         for (int j = 0; j < A.cols(); j++)
166             std::cout << A(i,j) << " ";
167     }
168     std::cout << "\n\n" << std::endl;
169     for (int i = 0; i < B.rows(); i++) {
170         std::cout << "\n" << std::endl;
171         for (int j = 0; j < B.cols(); j++)
172             std::cout << B(i,j) << " ";
173     }
174     std::cout << "\n\n" << std::endl;
175     for (int i = 0; i < C.rows(); i++) {
176         std::cout << "\n" << std::endl;
177         for (int j = 0; j < C.cols(); j++)
178             std::cout << C(i,j) << " ";
179     }
180     std::cout << "\n\n" << std::endl;
181 }

```

## D Matrix.h

```
1  #ifndef MATRIX_H_
2  #define MATRIX_H_
3
4  #include <vector>
5  #include <complex>
6
7  using std::vector;
8  using std::complex;
9
10 // simple wrapper of a vector
11 template <class T>
12 class Matrix {
13 public:
14     // create an empty matrix
15     Matrix(int numRows, int numcols)
16         :Nrow(numRows), Ncol(numcols), elements(Nrow*Ncol) {}
17
18     // construct it from existing data
19     Matrix(int numRows, int numcols, T* data)
20         :Nrow(numRows), Ncol(numcols), elements(data, data+numRows*numcols) {}
21
22     int rows() {return Nrow;}
23     int cols() {return Ncol;}
24
25     // access to elements, col is the fast axis
26     T operator() (int row, int col) const {return elements[Ncol*row + col];}
27     T& operator() (int row, int col) {return elements[Ncol*row + col];}
28
29
30     // get raw pointer to elements[0]
31     T* data() {return elements.data();}
32     const vector<T>& elem() {return elements;}
33
34 private:
35     int Nrow, Ncol;
36     vector<T> elements;
37 };
38
39
40 #endif // MATRIX_H_
```

## References

- [1] Introduction to Parallel Computing.” Second Edition - Grama, Gupta, et al. - 2003 Available online here:  
<http://parallelcomp.uw.hu/ch03lev1sec1.html>
- [2] <http://www.acms.arizona.edu/CompResources/index.html>
- [3] AFIPS '67 (Spring) Proceedings of the April 18-20, 1967, spring joint computer conference Pages 483-485